

Presentation writer: Tigran Hayrapetyan

Lecturer | Programmer | Researcher

[www.linkedin.com/in/tigran-hayrapetyan-cs/](http://www.linkedin.com/in/tigran-hayrapetyan-cs/)

## K-d trees

*prerequisites:*

Binary Search Tree.

# K-d trees

## *introduction*

There are multiple problems related to **multidimensional searches**:


- range search,
- nearest neighbor search,
- handling extensional objects, ...

And there is a variety of data structures which address that problems:

- Quadtrees,
- K-d trees,
- Range trees,
- Interval trees, ...

However, more often K-d trees are used for **nearest neighbor searches**, as they provide better time complexity there.

# Nearest neighbor search



Given a collection of points in  $N$ -dimensional space,

And a query point in it,

Find the point which is **closest** to the query point.

*Definition of NN search problem*

# Nearest neighbor search

*usage*

## 1) Various geometrical problems:

*approximate search for objects...*

## 2) Machine learning:

*supervised ML, both classification and regression (better known as **k-NN**, standing for k-nearest neighbors)...*

## 3) N-body problem:


*simulating gravitational interaction between  $N$  bodies, precalculating influence of entire regions to individual bodies...*

## 4) Color reduction:

*compress given image, using at most  $M$  colors to represent it...*

# Nearest neighbor search

*trivial solutions*



Sorting all points only by  
X-coordinate will not help...

Closest by X-coordinate doesn't  
mean the **closest**.

... and vice versa.

*Sorting all points only by X-coordinate*

# Nearest neighbor search

*trivial solutions*

Same about sorting only by Y-coordinate, or any other direction...

So, we must somehow take into account both X and Y coordinates **simultaneously**.

# K-d trees

Invented in **1975**, by **Jon L. Bentley**.

Jon Louis Bentley,  
Stanford University



# K-d trees

At first, let's concentrate on **2-dimensional** K-d trees:



Structure:

- 1) A simple binary tree.
- 2) Every 2D point corresponds to a node, and vice-versa, both leaves and intermediate nodes do store points.
- 3) Intermediate nodes can act slightly different:  
an intermediate node splits its area either by **X-axis** (vertical line), or by **Y-axis** (horizontal line).



# K-d trees

*examples*




*An incremental example of a K-d tree.*

# K-d trees


*properties*

Generally, the axes by which splits are performed, do interleave:  
... so we have – **XYXYXY**...

**Property #1:** Area corresponding to each node is more like a **"square"** and not a **"rectangle"**.




vs



# K-d trees

## *properties*

**Property #2:** If K-d tree has " $n$ " points, then:  
*best* partitioning will result in " $\log_2 n$ " tree height,  
*worst* partitioning will result in " $n$ " tree height.



# K-d trees

*properties*

Those estimates **do not depend** on  
coordinates of the points, neither on  
sequence of split axes e.g. **XYXYXY...**



*Exercise:*

*In which order points of the previous example should be inserted, for us to have a more or less balanced K-d tree?*

# K-d trees

*properties*

**Property #3:** After any split, subtree of the first half is **not related** anymore to subtree of the second half.



# K-d trees

## *exercises*

*Exercise:*


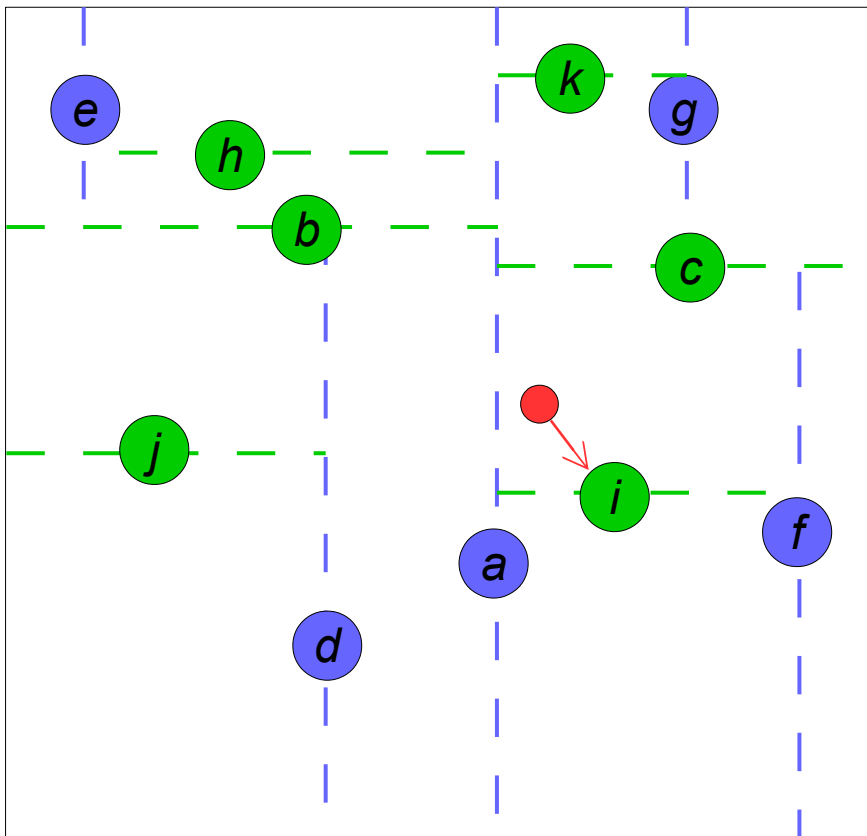
*Insert the following points into an initially empty K-d tree:  
(2,3) , (8,7) , (5,1) , (1,9) , (7,4) , (4,8) , (3,2) , (9,6).*

# K-d trees

*nearest neighbor search*

Top use of K-d trees is to perform **nearest neighbor** (NN) search.



Given a **query point** "q", find the point from K-d tree which is closest to "q".



# K-d trees

*nearest neighbor search*

We'll try to search by the correct branch in the tree, every time **updating current best** distance.







# K-d trees

*nearest neighbor search*

On this example we have found the proper answer - "i".

However, if there would be some extra **points in the left half**, our answer will become incorrect.





So, do we need to observe also the other branch?

# K-d trees

*nearest neighbor search*

If we decide to always observe also the other branch, then we should do it on every level...

... which will result in **inspecting entire tree**, and  $O(n)$  time complexity.





# K-d trees

*nearest neighbor search*

So how should we do then... ?

**Key idea** is: during exit from recursion, inspect the other half too only if distance to it is less than current best minimum.


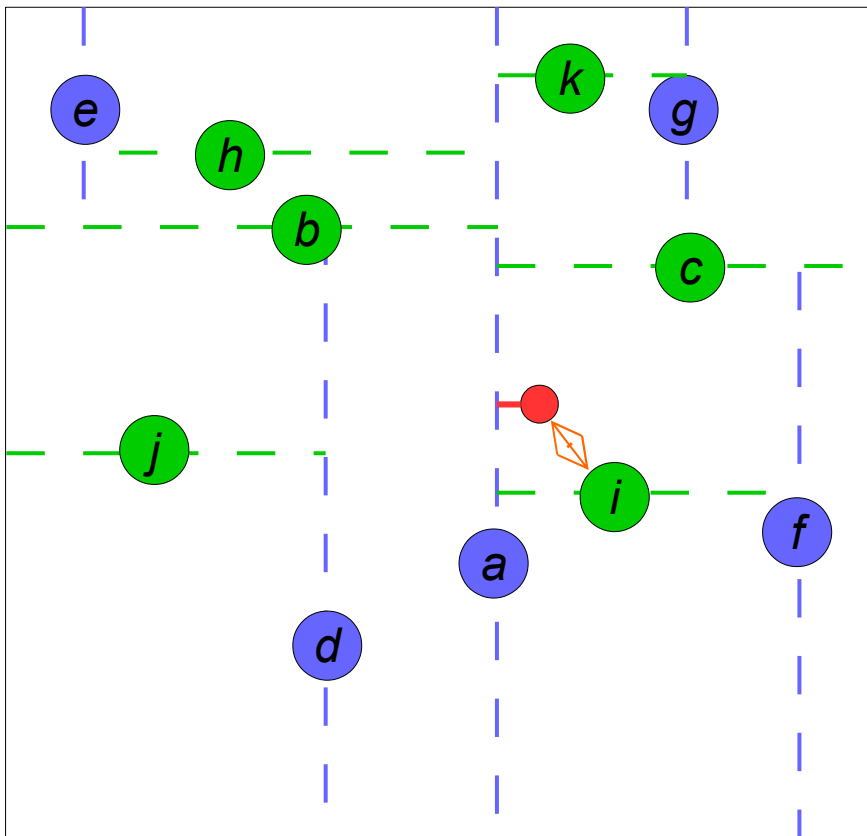


So, in this example, **only 1** extra branch need to be inspected.

# K-d trees

*nearest neighbor search*

Once we do that too, here are the edges of the tree which will be inspected at the end.



# K-d trees

## *nearest neighbor search*

Now we can write pseudocode of the nearest neighbor search procedure:

```
procedure NN_Search( n: Node, q: Point, currMin: NumberRef,  
                      currNN: NodeRef )  
    // Update current minimum  
    if dist(q,n) < currMin  
        currMin := dist(q,n)  
        currNN := n  
    // Branch the search  
    if q is on left (or bottom) half of n  
        NN_Search( left{n}, q, currMin, currNN )  
        // If need to check also the other branch  
        if dist(q,plane{n}) < currMin  
            NN_Search( right{n}, q, currMin, currNN )  
    else  
        NN_Search( right{n}, q, currMin, currNN )  
        // If need to check also the other branch  
        if dist(q,plane{n}) < currMin  
            NN_Search( left{n}, q, currMin, currNN )
```

# K-d trees

*nearest neighbor search*

We have designed the algorithm for NN search.

... but **how efficient** is it?

# K-d trees


*nearest neighbor search*

**Observation #1:** During the NN search:

current best minimum can  
only decrease,

while distance to the other half  
generally increases,

so it becomes **less and less probable**,  
that we will need to observe  
also the other half.





# K-d trees

*nearest neighbor search*

**Observation #2:** Need to observe also the other half-plane will generally occur only for such query points, which are "*close enough*" to the other half-plane.

And that holds generally for nodes which are **deep enough**.





# K-d trees

*nearest neighbor search*

So time complexity of NN search is only by **some constant** greater than tree height " $h$ ", so it gives:

$$O(h) = O(\log n)$$

if the points are **distributed uniformly**.

The NN search algorithm can be easily extended to find k-nearest neighbors – **k-NN**, for the query point " $q$ ". In order to do that:

we keep " $k$ " current best minimums, instead of just 1,


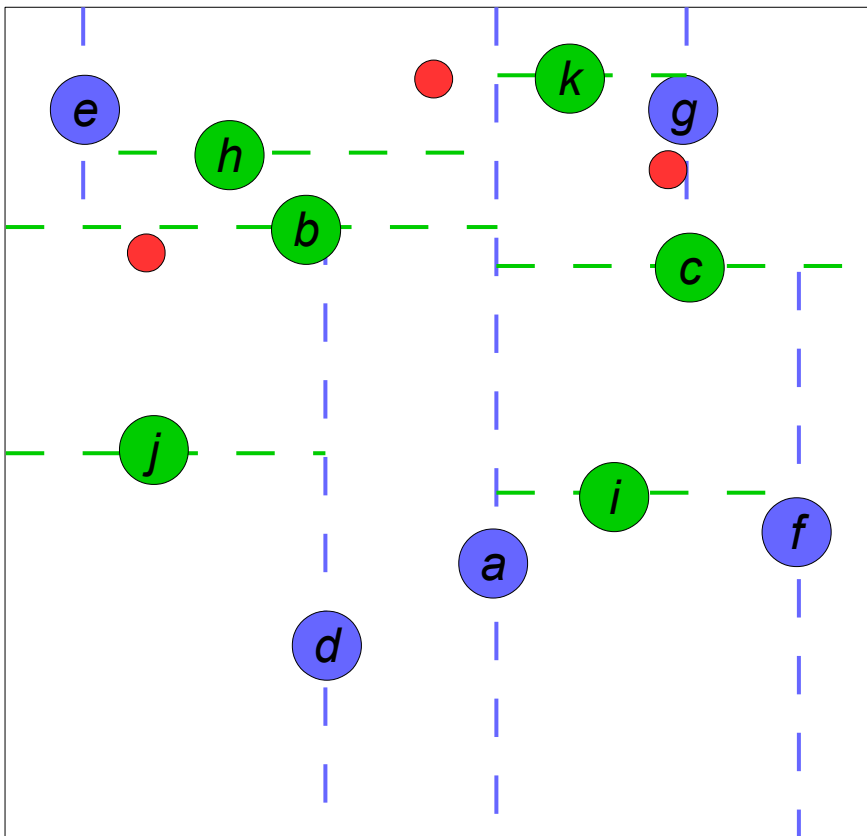
and we inspect also the other half-plane if distance to it is less than any of the " $k$ " current bests.

# K-d trees

*nearest neighbor search*

*Exercises:*


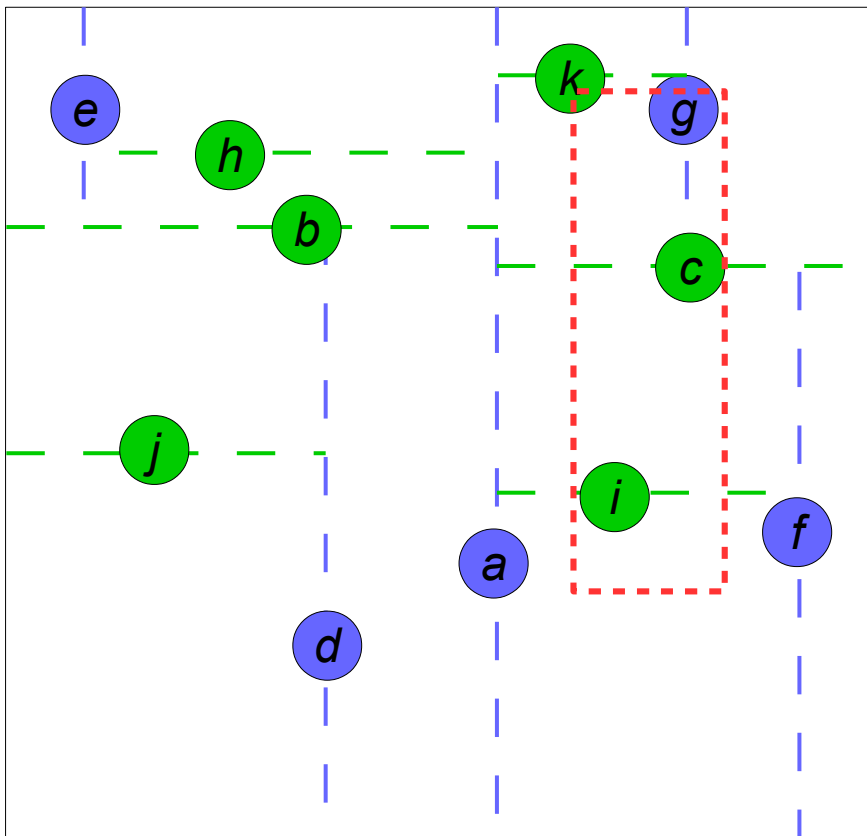
*For given K-d tree, perform NN search from the specified points:*



# K-d trees

*range search*

For a range search we need to report all the points which are inside **query rectangle "q"**.




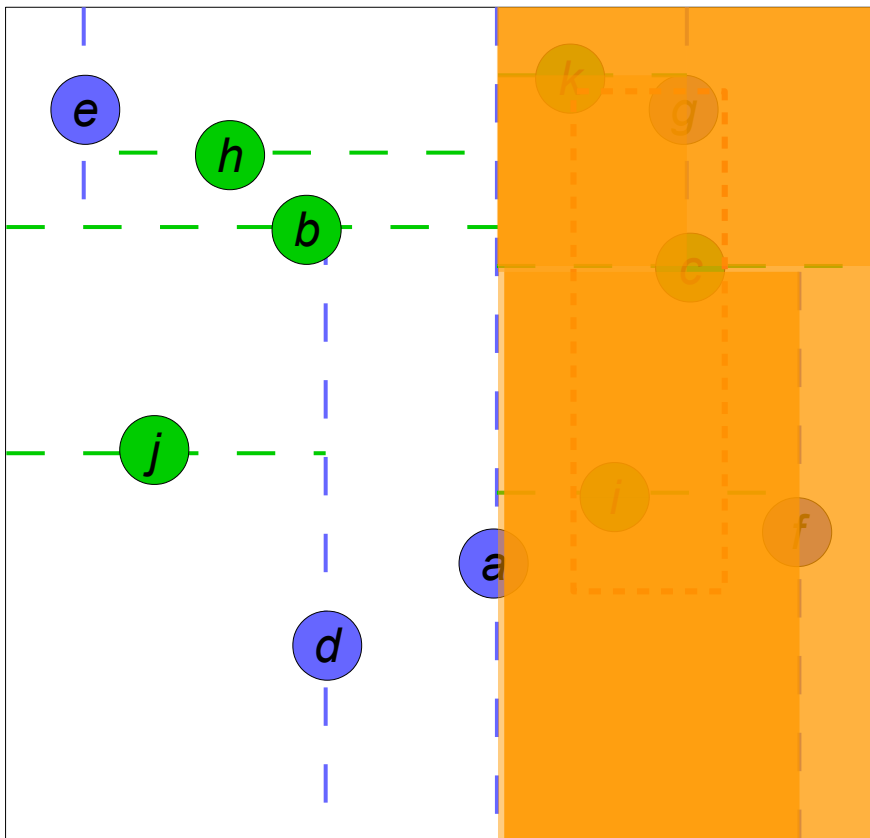
So, here we need to report  
"g", "c", "i".

# K-d trees

*range search*

The algorithm is simple:

if "q" fits entirely in only half-plane, we continue from there,  
otherwise we continue from both half-planes.





# K-d trees

*range search*

But what is **time complexity** of such range search?

At first, let's consider the case when " $q$ " is very **thin** and very **long**.



# K-d trees

*range search*

As we can see:

the search splits on every **odd** level,  
and the search doesn't split on every **even** level.

So we have:

$f(n) = 2 * f(n/2)$  for odd levels,

$f(n) = f(n/2)$  for even levels,

or combining it:

$f(n) = 2 * f(n/4)$

Continuing recursively, we obtain:

$$f(n) = 2 f\left(\frac{n}{4}\right) = 4 f\left(\frac{n}{16}\right) = 2^{\frac{k}{2}} f\left(\frac{n}{2^k}\right) = 2^{\frac{\log_2 n}{2}} = 2^{\log_2 n * \frac{1}{2}} = \sqrt{2^{\log_2 n}} = \sqrt{n}$$


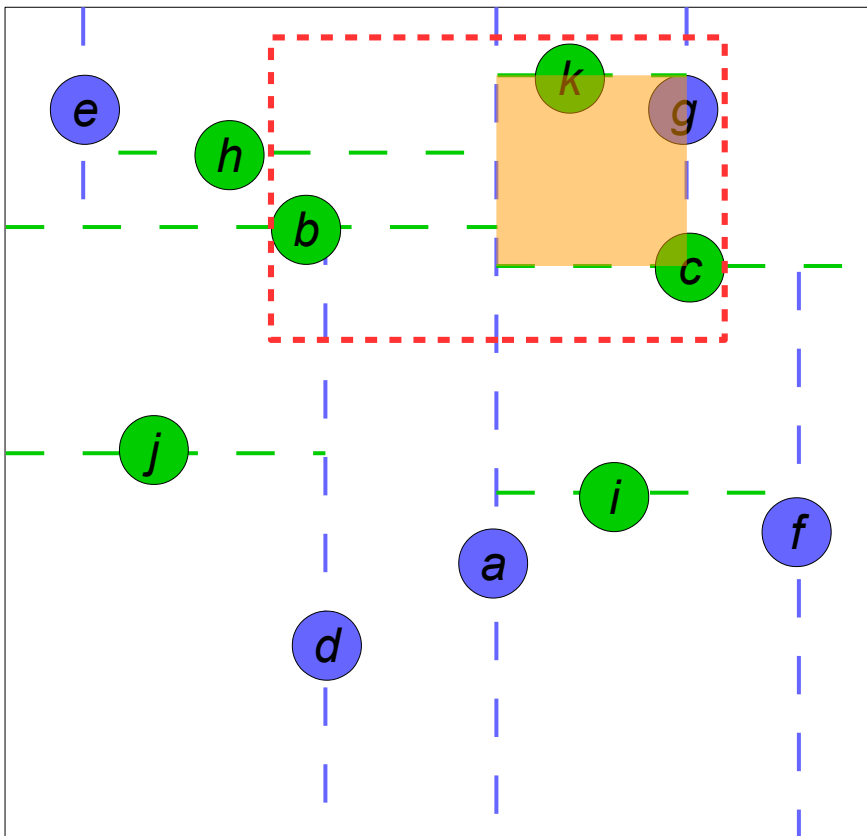
So time complexity for such thin rectangle is  $O(\sqrt{n})$ .

# K-d trees

*range search*

But what about a **normal** query rectangle?

Let's note that for some "q" there can be such subtrees, which will be **reported completely**. That is when their area lies inside "q".




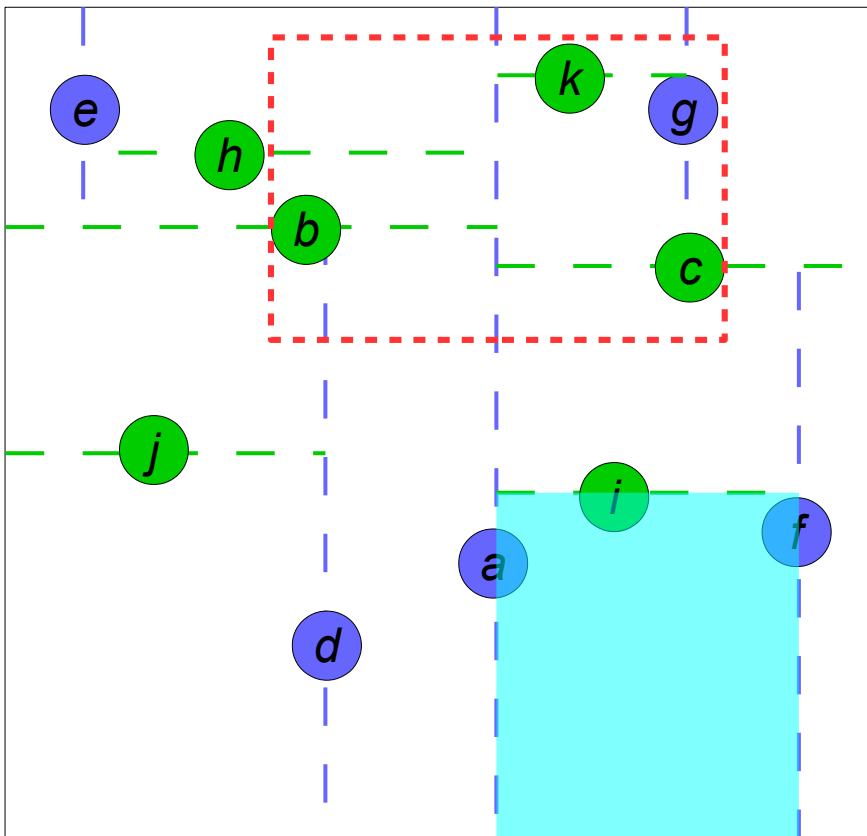
Here, subtree left to "k"  
will be reported completely.

# K-d trees

*range search*

From the other side, some other subtrees will **not be touched** at all.

Here, subtree left to "i" will not be touched.



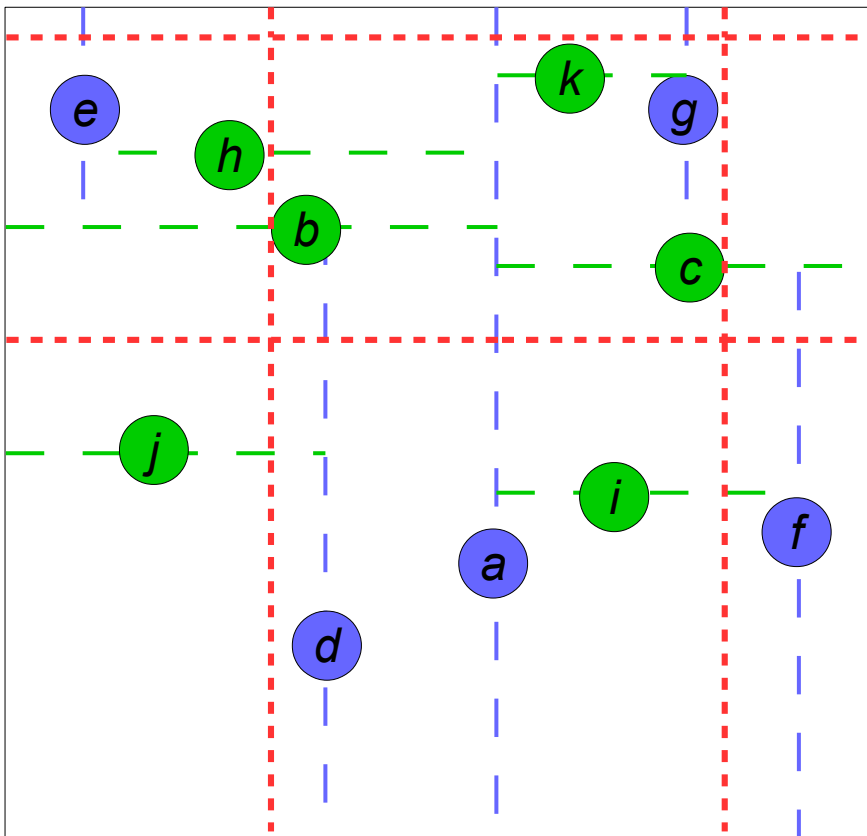


# K-d trees

*range search*

It remains to analyze the subtrees, which will be **reported partially**.

And number of such subtrees is not greater than of those, intersecting the following 4 thin rectangles.



So we have:

Completely reported subtrees:  
 $O(k)$ ,

Partially reported subtrees:  
 $O(4 \cdot \sqrt{n}) = O(\sqrt{n})$ ,

Not touched subtrees:  
no time spent.

Which in sum gives:


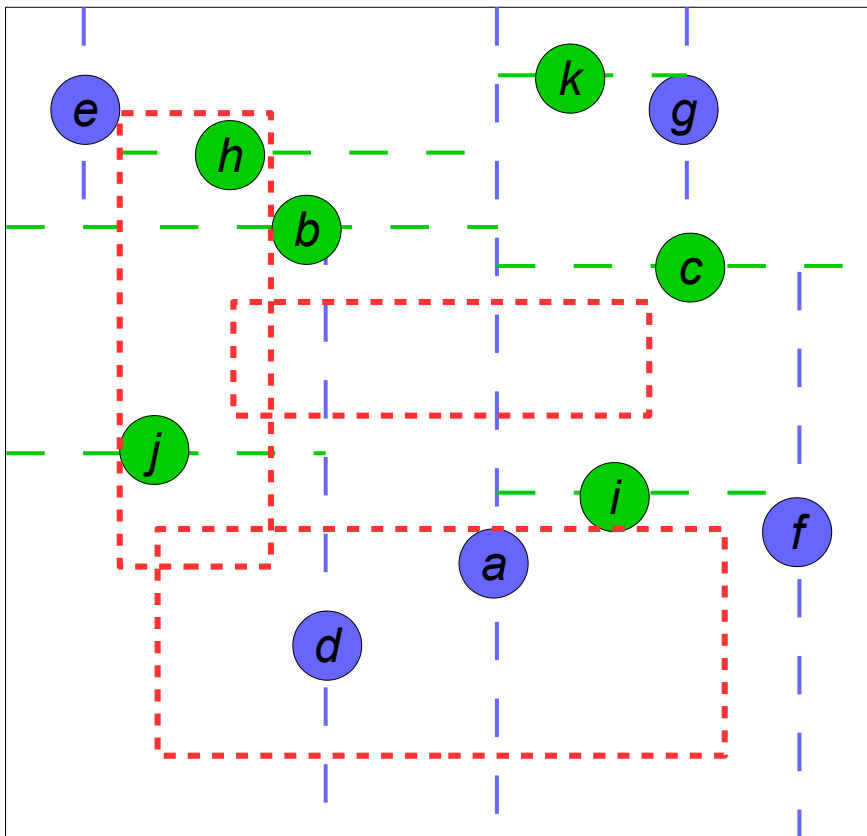
$O(\sqrt{n} + k)$ , where "k" is number of reported points.

# K-d trees

*range search*

*Exercises:*


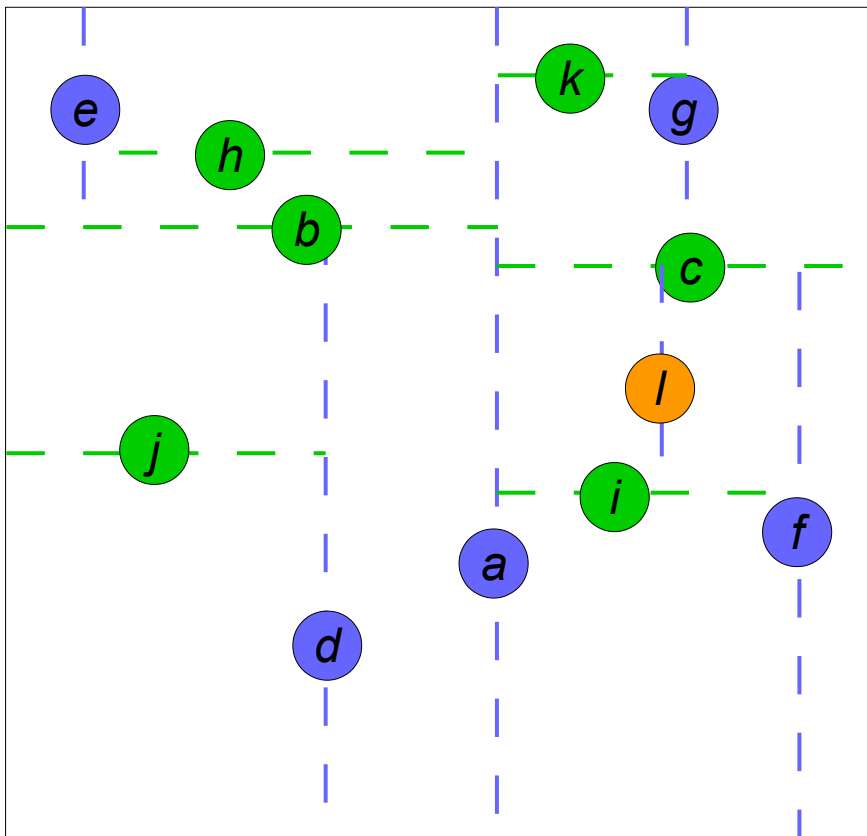
*Perform range searches with the following query rectangles:*



# K-d trees

*adding points*

Adding new point is quite similar to checking for presence:


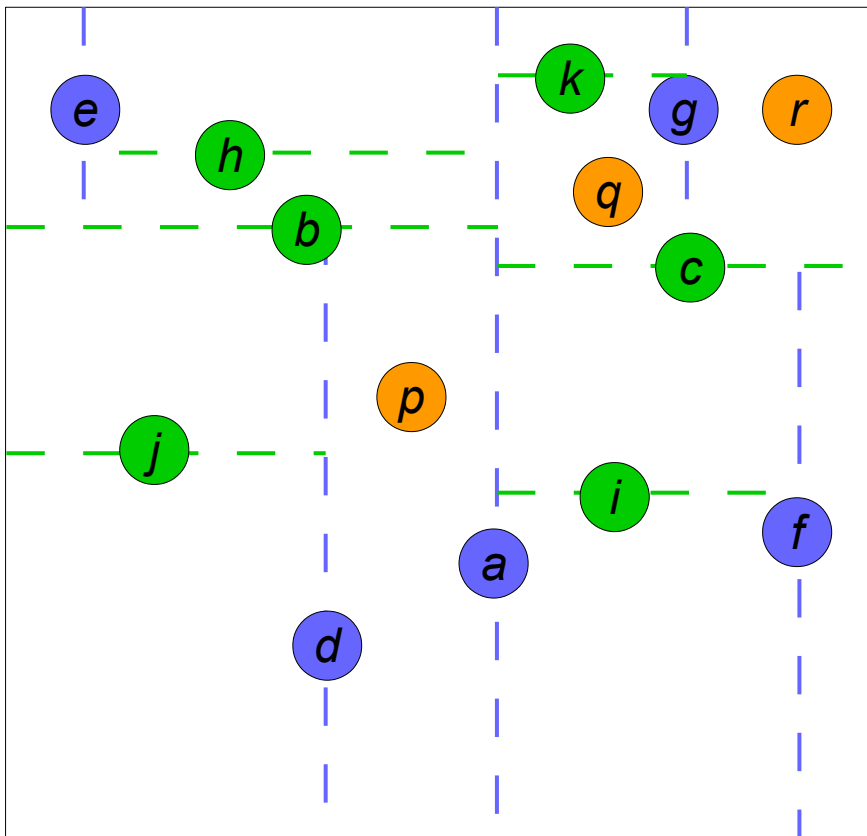


# K-d trees

*adding points*

Exercises:


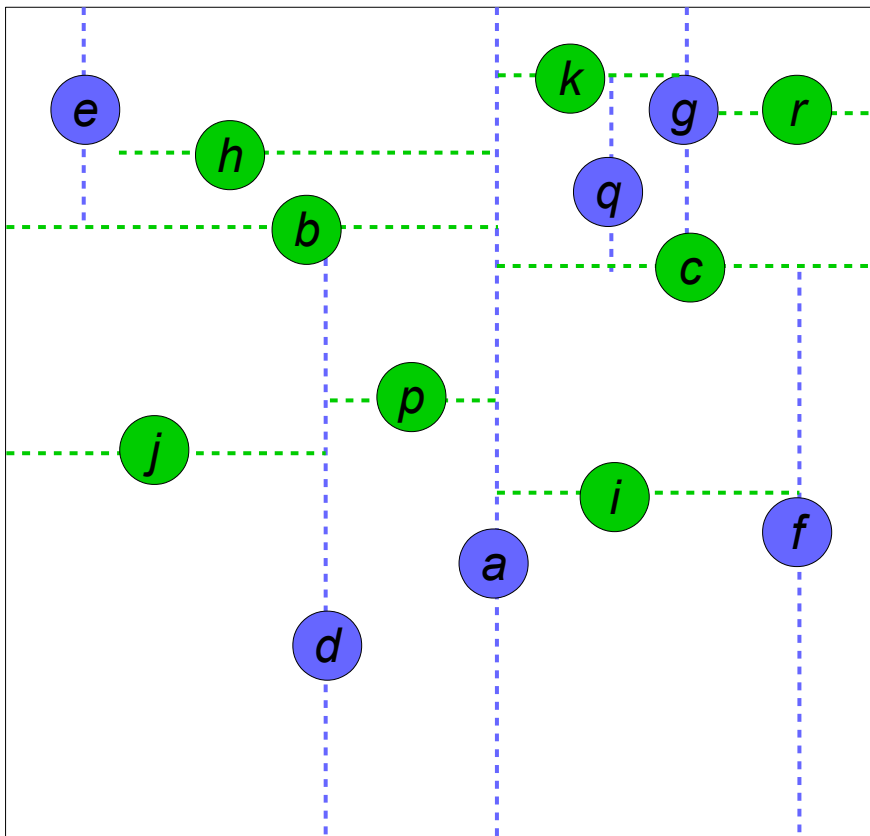
Insert the following points into the K-d tree:



# K-d trees

*removing points*



If node is a leaf, removal is trivial.



# K-d trees

*removing points*

Other scenario: if it has descendants, it must be replaced by the "nearest" one from them (i.e. nearest one from it's region).





# K-d trees

*removing points*

But it can be also that the "nearest" one has descendants too.

In that case it must be removed from the tree recursively.


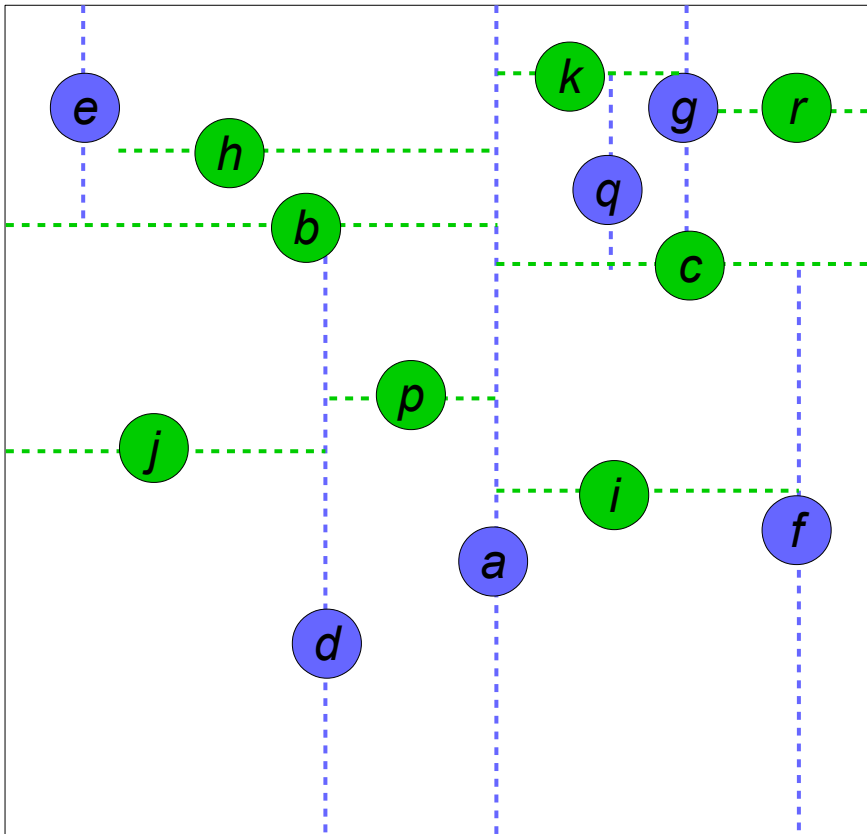


# K-d trees

*removing points*

## Exercise 1:

Remove the following point from the K-d tree:




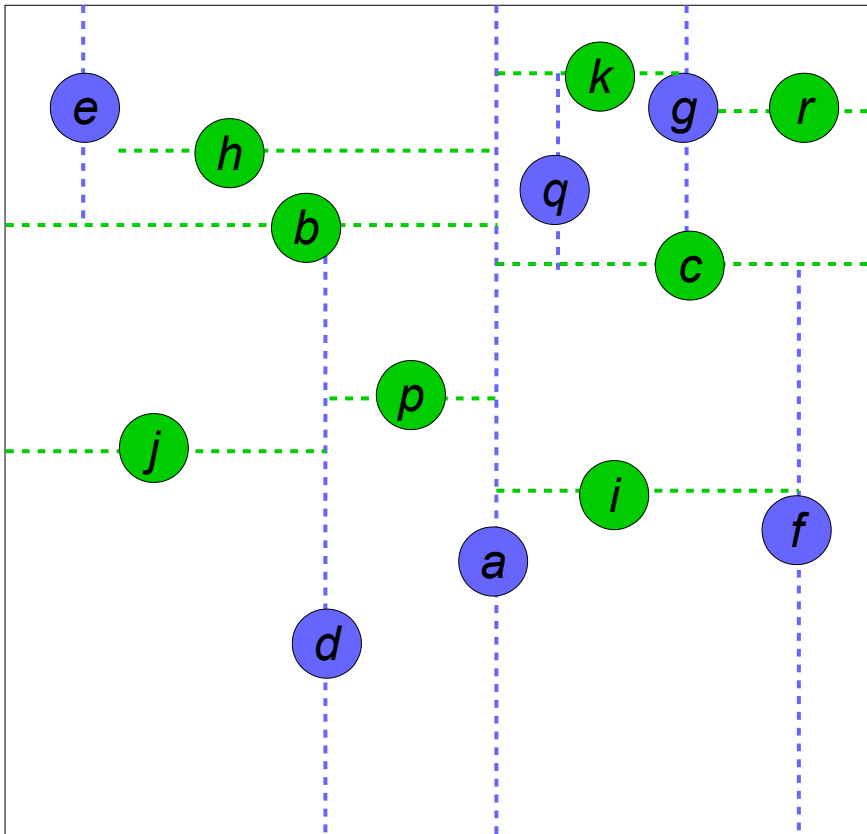


# K-d trees

*removing points*

## Exercise 2:

Remove the following point from the K-d tree:





# K-d trees

*batch construction (approach 1)*

The most straightforward way:

- 1) Find the median,
- 2) Partition by it,
- 3) Recursively construct subtrees.



# K-d trees

*batch construction (approach 1)*


So the question is: **how to find median points?**

<i>method</i>	<i>result</i>	<i>time</i>
Median of a sample	not always a good median	$O(1)$
Complete sort	best result	$O(N \cdot \log N)$
QuickSort-based median	best result	expected $O(N)$

# K-d trees

*batch construction (approach 2)*


Sort the points **by both X and Y** coordinates. So at each step:  
find the median & divide into 2 parts,  
rewrite 1<sup>st</sup> items & 2<sup>nd</sup> items from the other sequence,  
we again have necessary sorted lists: continue recursively.



# K-d trees

*batch construction (approach 2)*

Having that said, on each step we have both sorted lists.



Picking median is trivial.

Time to construct 2 pair of smaller lists, from lists of size N is:

$O(N)$ .

Overall time complexity of batch construction:

$O(N \log N)$  for sorting by both X and Y.

$N + 2 \cdot (N/2) + 4 \cdot (N/4) + \dots + N \cdot 1 = O(N \log N)$  for construction itself.

Summary:  **$O(N \log N)$** .



# K-d trees

*higher dimensions*

K-d trees can be easily extended to **3D** and higher dimensions.

All what we should do is just to perform splits by multiple axes:

*XYZXYZXYZ...*



Generally, all algorithms of the K-d tree **remain unchanged**.

# K-d trees

*higher dimensions*

Time complexity of range search will be:

$$O\left(D * n^{1-\frac{1}{D}} + k\right) \quad \text{where } D \text{ is number of dimensions.}$$

We can note that together with increase of  $D$ , time complexity of range search becomes more and more **closer to linear**.

In some sense, this holds also for NN search. For example, if  $n \sim D$  then NN search runs almost with same performance, as linear search.

If we want K-d tree to behave efficiently, we must ensure that  $n \ll D$ .

# K-d trees

*comparison with other data structures*

	<i>Quadtree</i>	<i>K-d tree</i>	<i>Range tree</i>
<i>NN Search:</i>		$O(\log N)$	
<i>Range search:</i>		$O(DN^{1-1/D} + k)$	$O(\log^D N + k)$
<i>Construction:</i>	$O(DN \cdot \log N)$	$O(DN \cdot \log N)$	$O(N \cdot \log^{D-1} N)$
<i>Insert:</i>	$O(\log N)$	$O(\log N)$	$O(\log^D N)$
<i>Remove:</i>	$O(\log N)$	$O(\log N)$	$O(\log^D N)$
<i>Space:</i>	$O(2^D \cdot N)$	$O(N)$	$O(N \cdot \log^{D-1} N)$
<i>Higher dimensions:</i>	bad scaling	good scaling	good scaling